



Python API for Spark

Josh Rosen

# What is Spark?



**Fast and expressive cluster computing system**

**Compatible with Hadoop-supported file systems  
and data formats (HDFS, S3, SequenceFile, ...)**

Improves *efficiency* through in-memory computing primitives and general computation graphs

As much as 30x faster

Improves *usability* through rich APIs in Scala, Python, and Java, and an interactive shell

Often 2-10x less code

# RDDs

*Resilient Distributed Datasets*

Immutable, partitioned  
collections of objects

# Transformations

map  
filter  
groupBy  
join  
...

# Actions

count  
collect  
save  
...

# Example: Log Mining

```
val lines = spark.textFile("hdfs://...")
val errors = lines.filter(_.startsWith("ERROR"))
val messages = errors.map(_.split('\t')(2))

messages.filter(_.contains("foo")).count
```



# What is PySpark?

# PySpark at a Glance


Write Spark jobs  
in Python



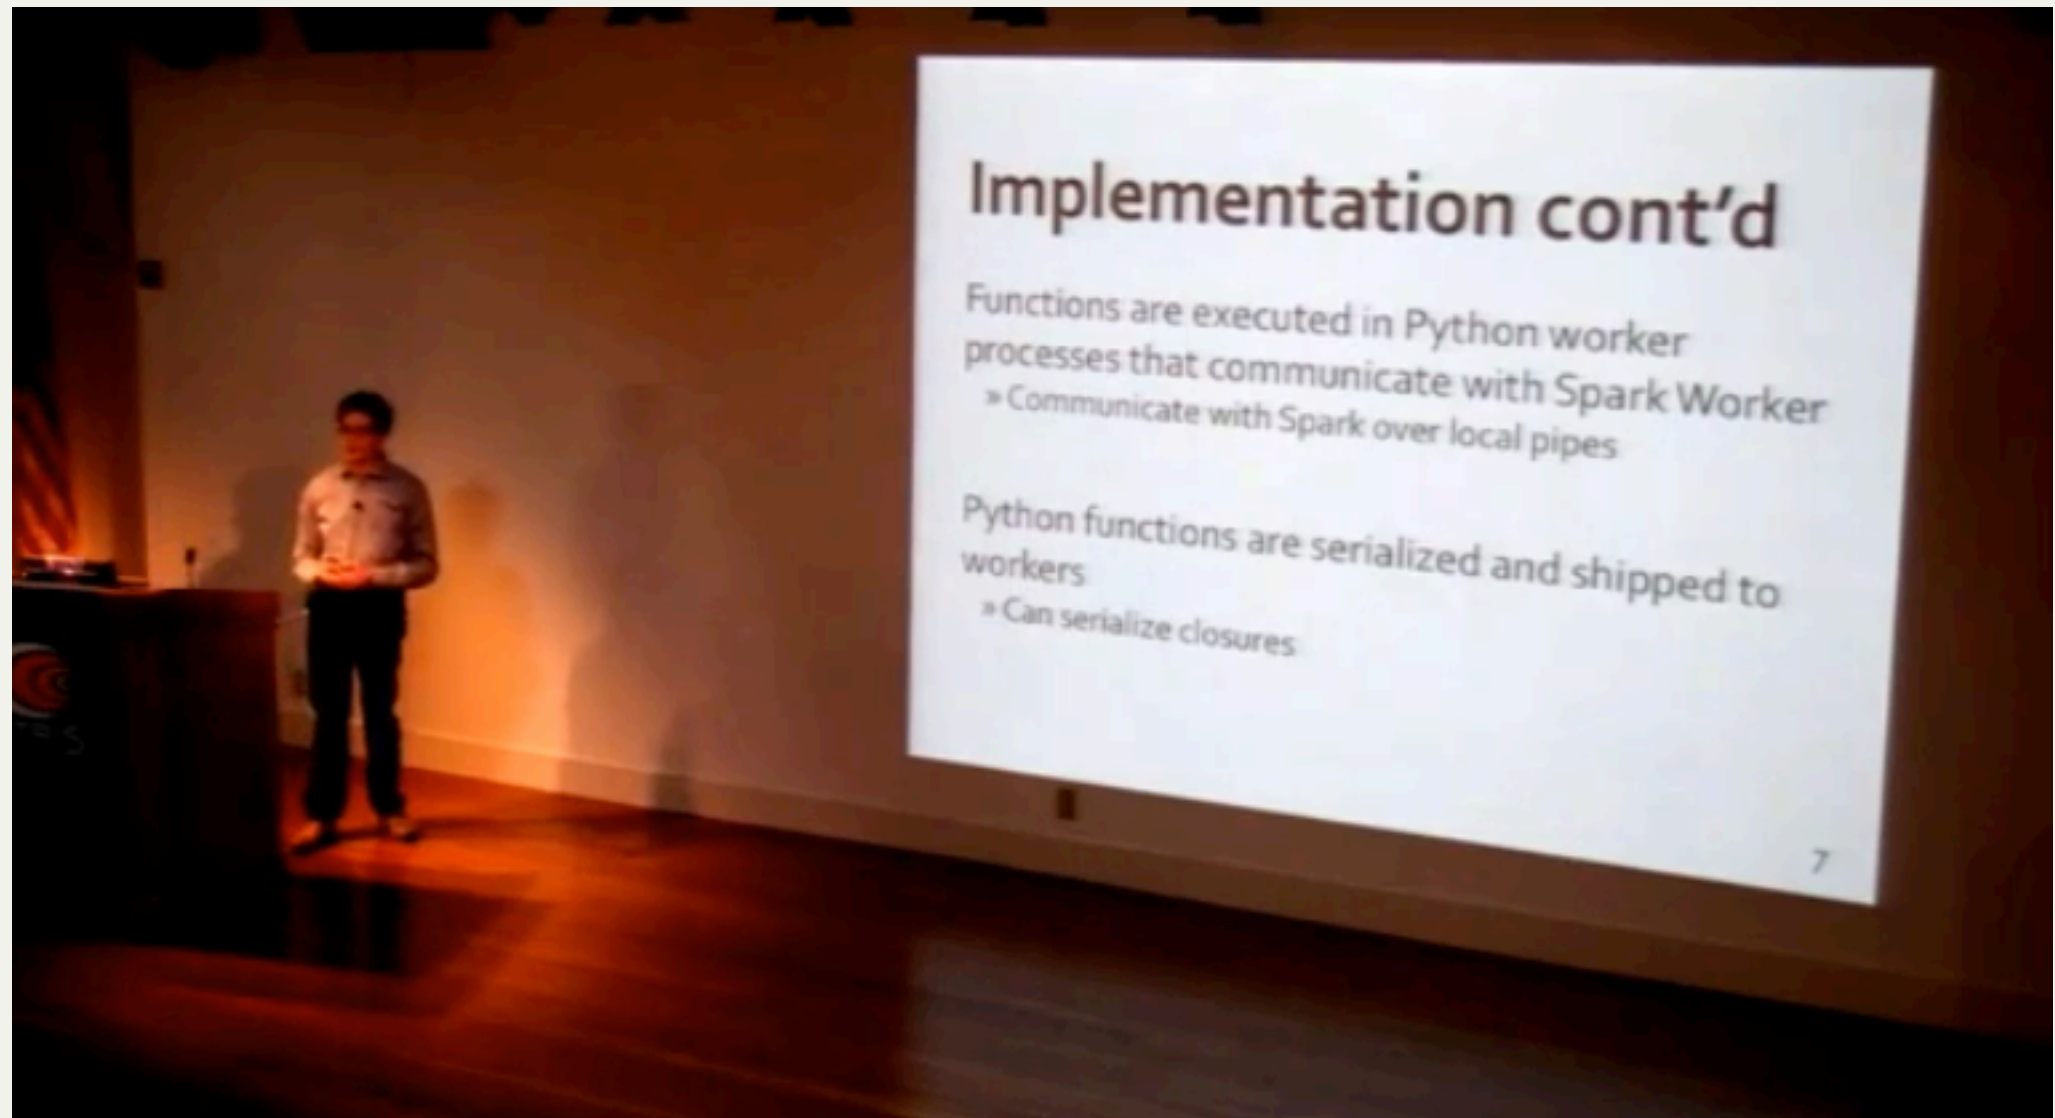
Run interactive  
jobs in the shell

```
2. Python
[joshrosen spark (master)]$ ./pyspark
Python 2.7.3 (v2.7.3:70274d53c1dd, Apr  9 2012, 20:52:43)
[GCC 4.2.1 (Apple Inc. build 5666) (dot 3)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
Spark context available as sc.
>>> sc.parallelize(range(1, 100)).map(lambda x: x**2).collect()
[1, 4, 9, 16, 25, 36, 49, 64, 81, 100, 121, 144, 169, 196, 225, 256, 289, 324, 361, 400,
, 441, 484, 529, 576, 625, 676, 729, 784, 841, 900, 961, 1024, 1089, 1156, 1225, 1296,
1369, 1444, 1521, 1600, 1681, 1764, 1849, 1936, 2025, 2116, 2209, 2304, 2401, 2500, 260
1, 2704, 2809, 2916, 3025, 3136, 3249, 3364, 3481, 3600, 3721, 3844, 3969, 4096, 4225,
4356, 4489, 4624, 4761, 4900, 5041, 5184, 5329, 5476, 5625, 5776, 5929, 6084, 6241, 640
0, 6561, 6724, 6889, 7056, 7225, 7396, 7569, 7744, 7921, 8100, 8281, 8464, 8649, 8836,
9025, 9216, 9409, 9604, 9801]
>>> 
```

Supports C  
extensions



# Previewed at AMP Camp 2012



# Available now in 0.7 release

# Example: Word Count

```
from pyspark.context import SparkContext

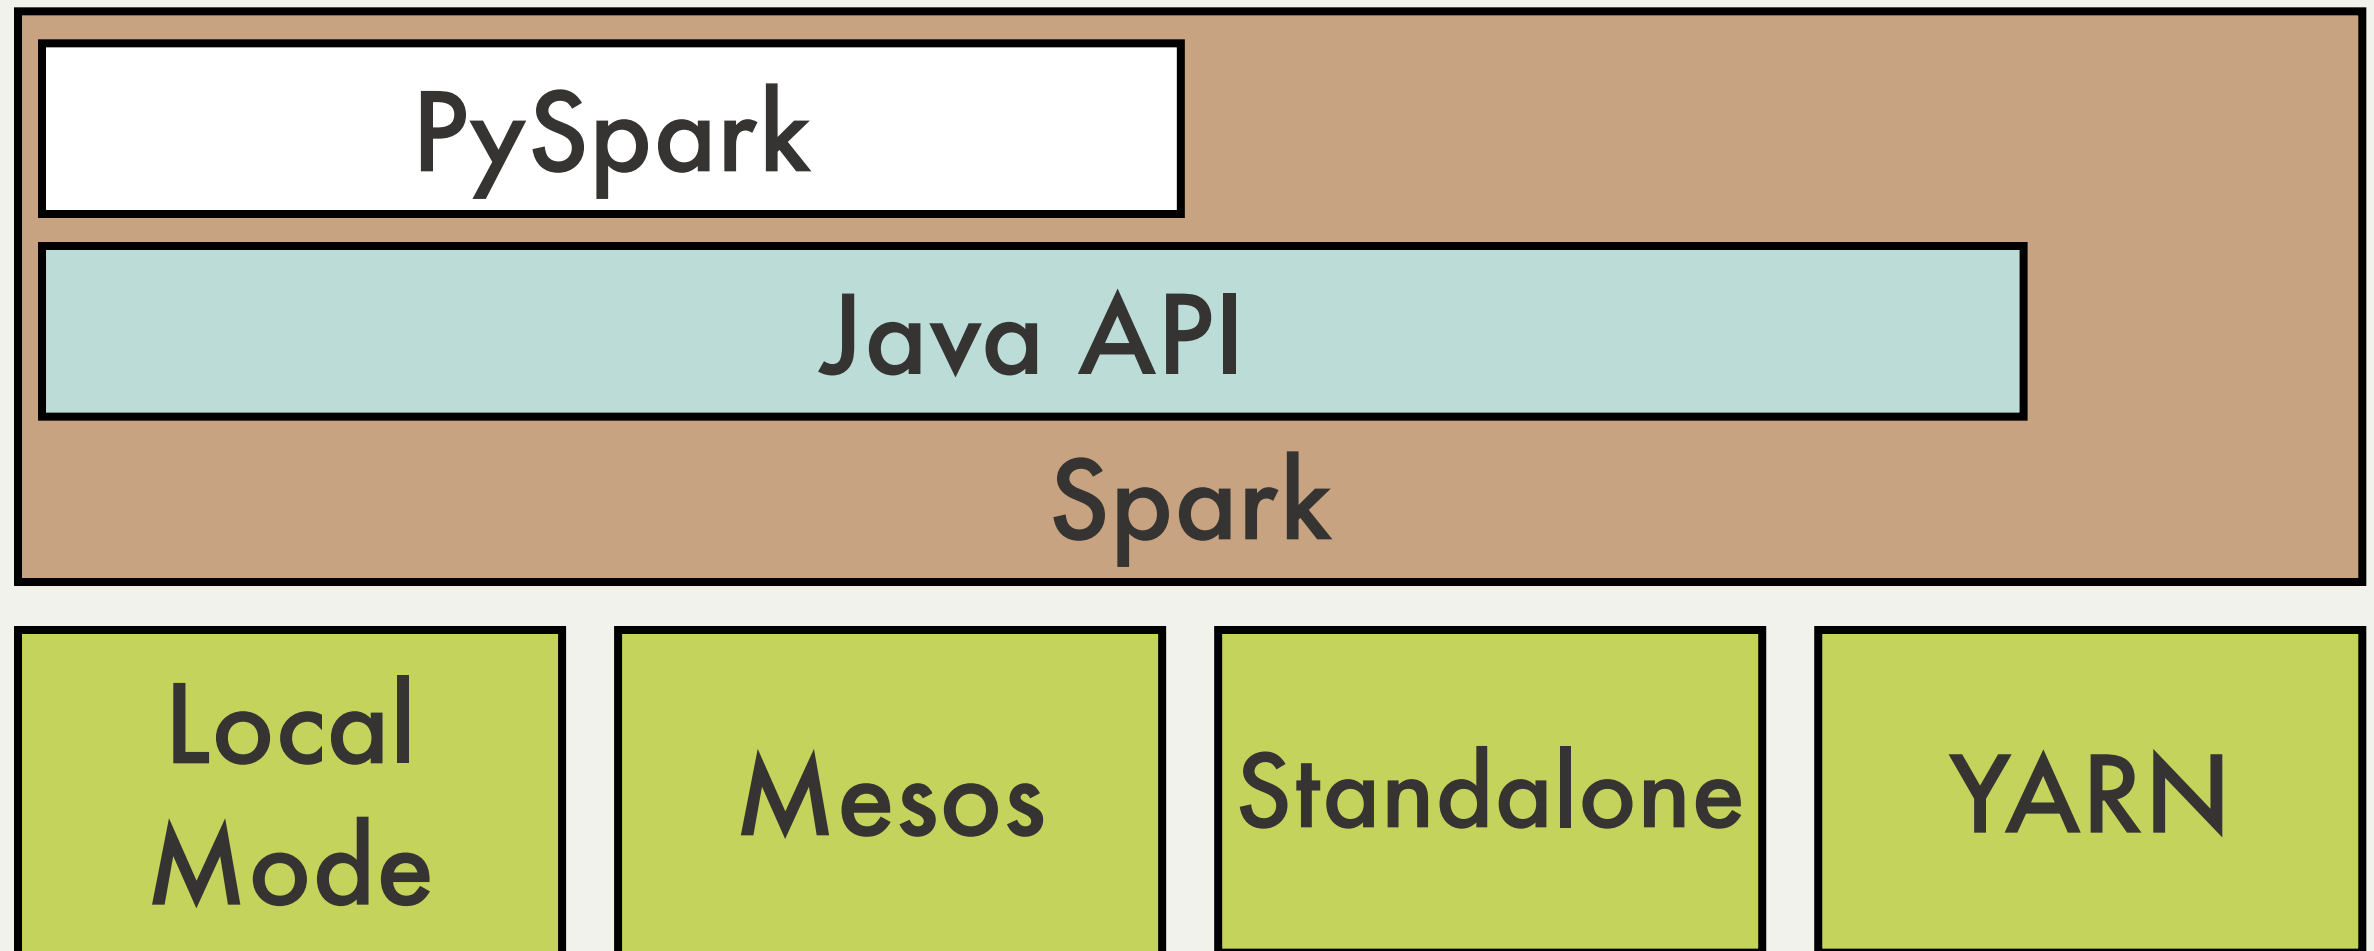
sc = SparkContext(...)
lines = sc.textFile(sys.argv[2], 1)
counts = lines.flatMap(lambda x: x.split(' ')) \
               .map(lambda x: (x, 1)) \
               .reduceByKey(lambda x, y: x + y)

for (word, count) in counts.collect():
    print "%s : %i" % (word, count)
```


Demo

# Implementation

# Built on top of Java API



# Process data in Python



and persist /  
transfer it in Java



**Re-uses Spark's** scheduling  
broadcast  
checkpointing  
networking  
fault-recovery  
HDFS access

# PySpark has a small codebase:

File	blank	comment	code
python/pyspark/rdd.py	115	345	302
core/src/main/scala/spark/api/python/PythonRDD.scala	33	45	231
python/pyspark/context.py	32	101	133
python/pyspark/tests.py	26	11	84
python/pyspark/accumulators.py	37	91	70
python/pyspark/serializers.py	21	7	55
python/pyspark/join.py	15	27	50
python/pyspark/worker.py	8	7	44
core/src/main/scala/spark/api/python/PythonPartitioner.scala	5	9	34
pyspark	9	8	27
python/pyspark/java_gateway.py	5	7	26
python/pyspark/files.py	7	14	17
python/pyspark/broadcast.py	8	16	15
python/pyspark/shell.py	4	6	8
python/pyspark/__init__.py	6	14	7
SUM:	331	708	1103

< 2K lines, including comments

# Data Flow

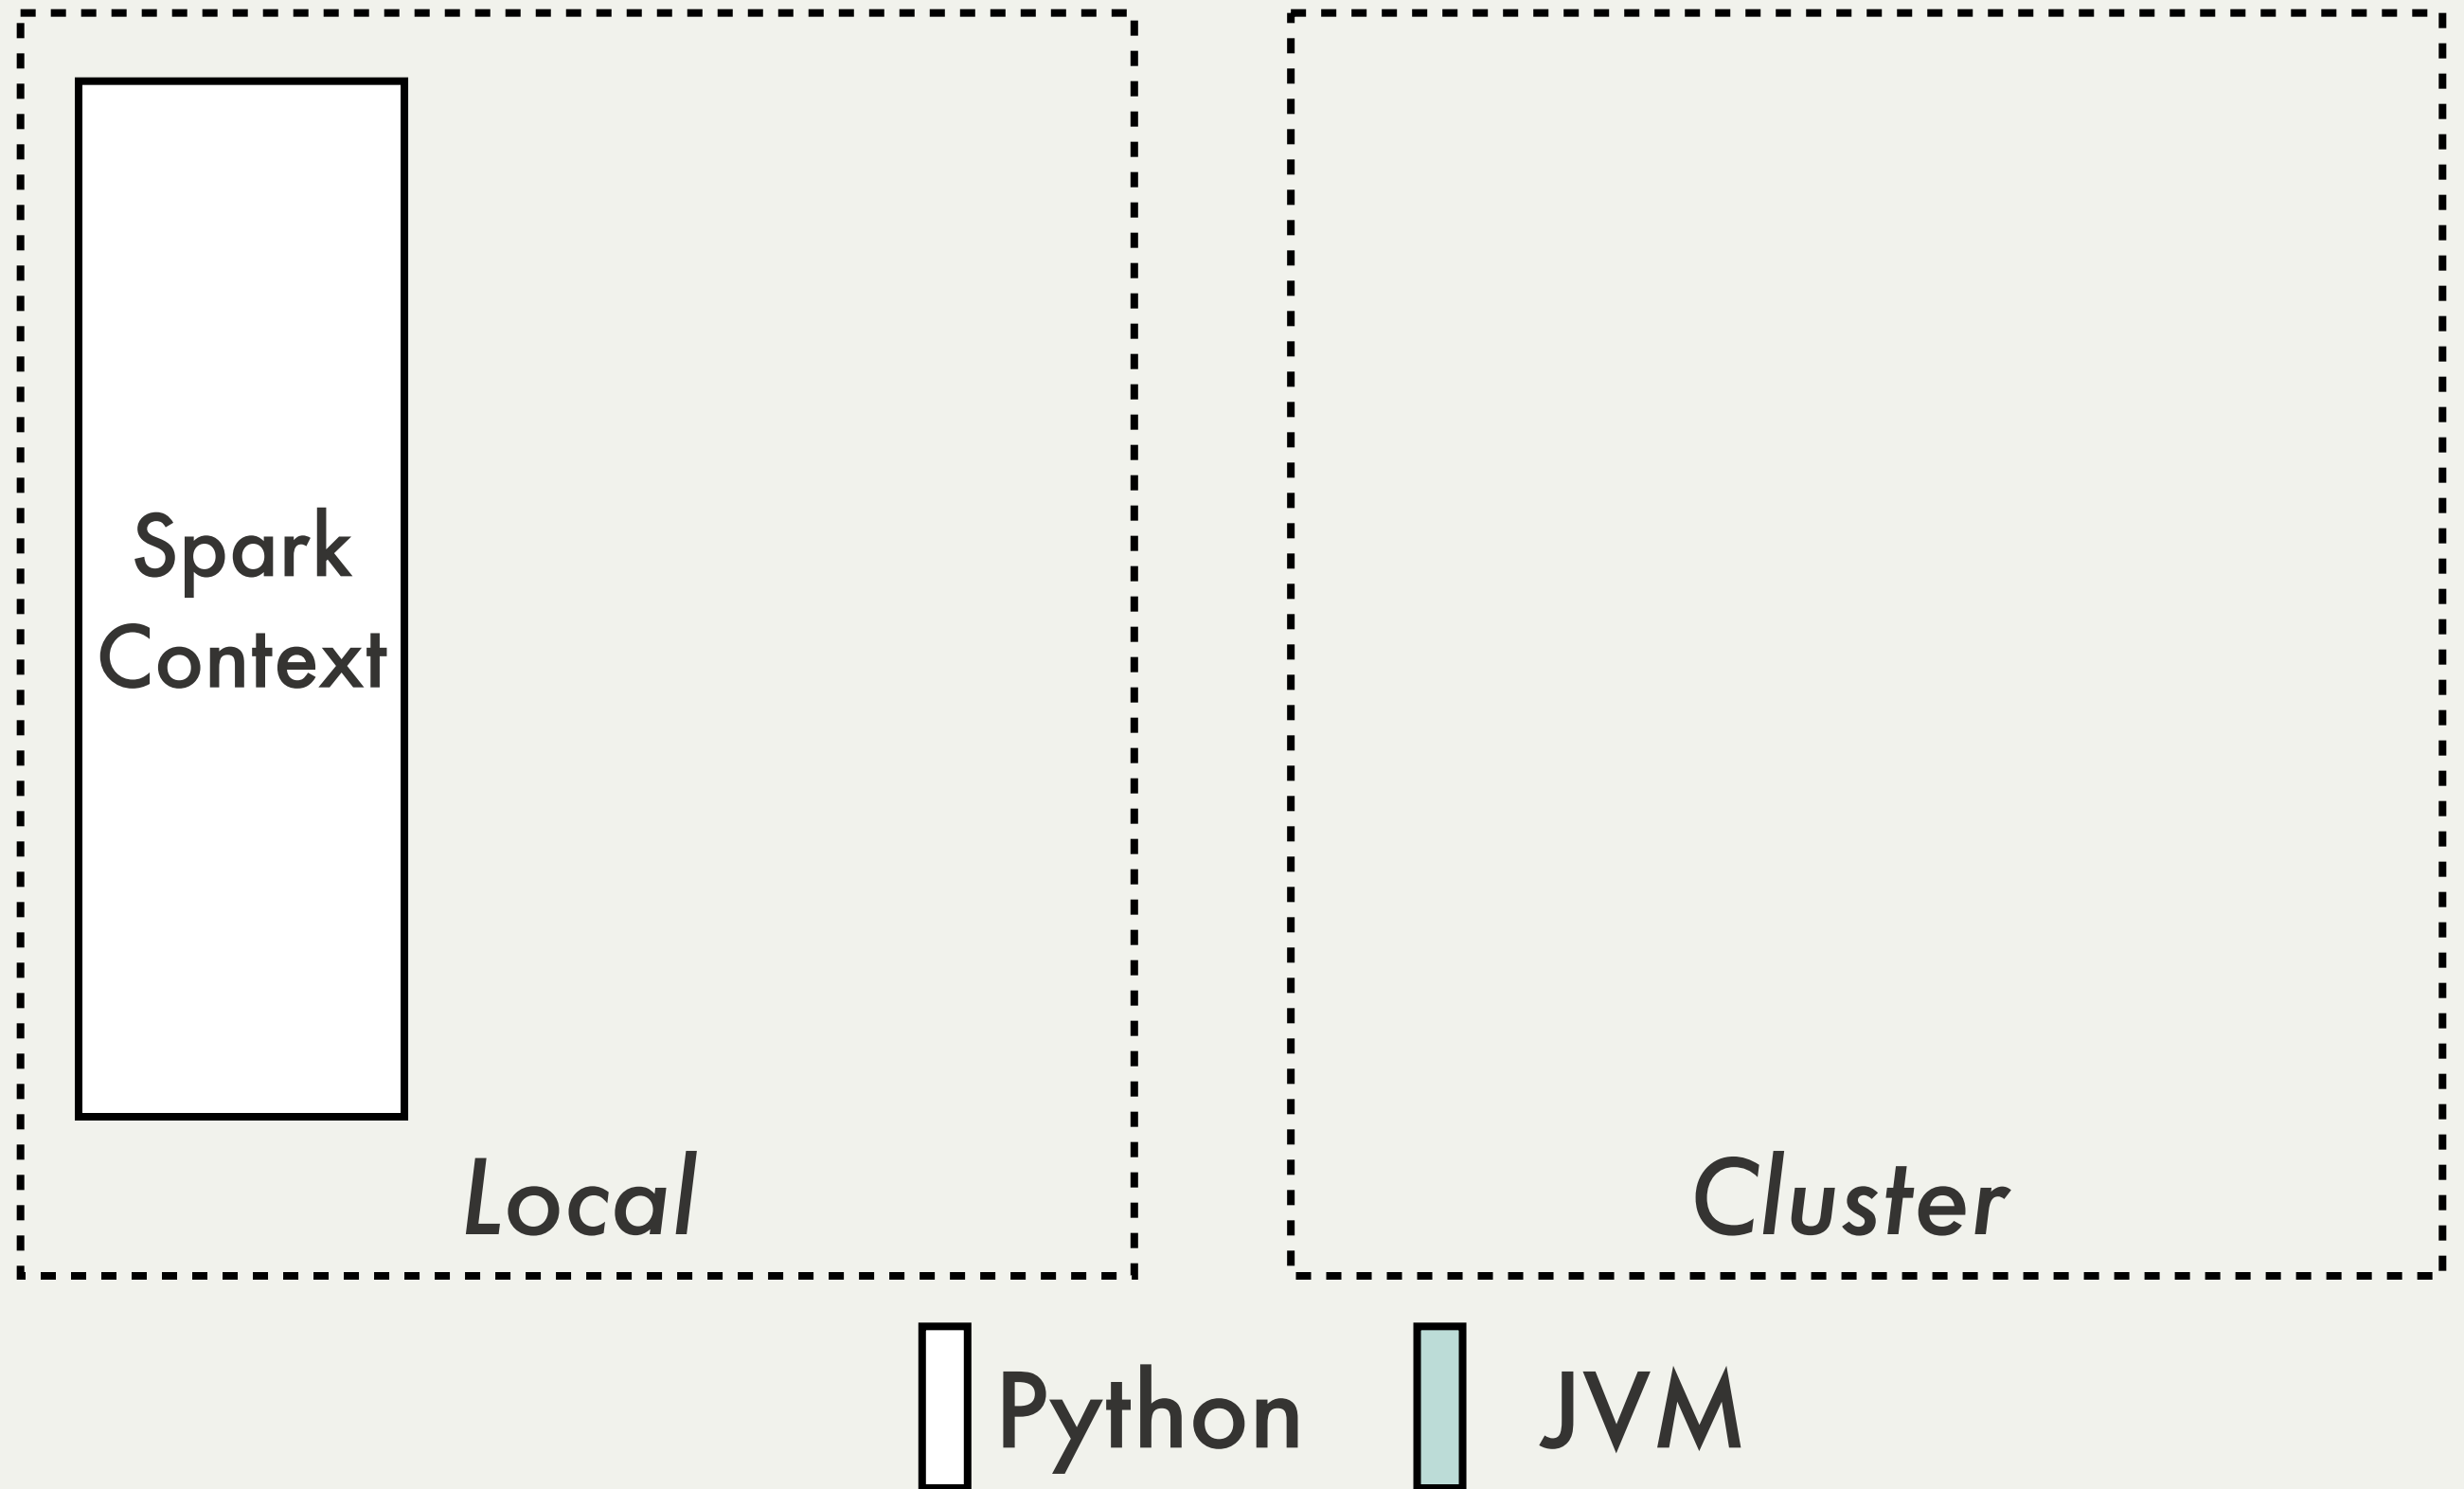
*Local*

*Cluster*

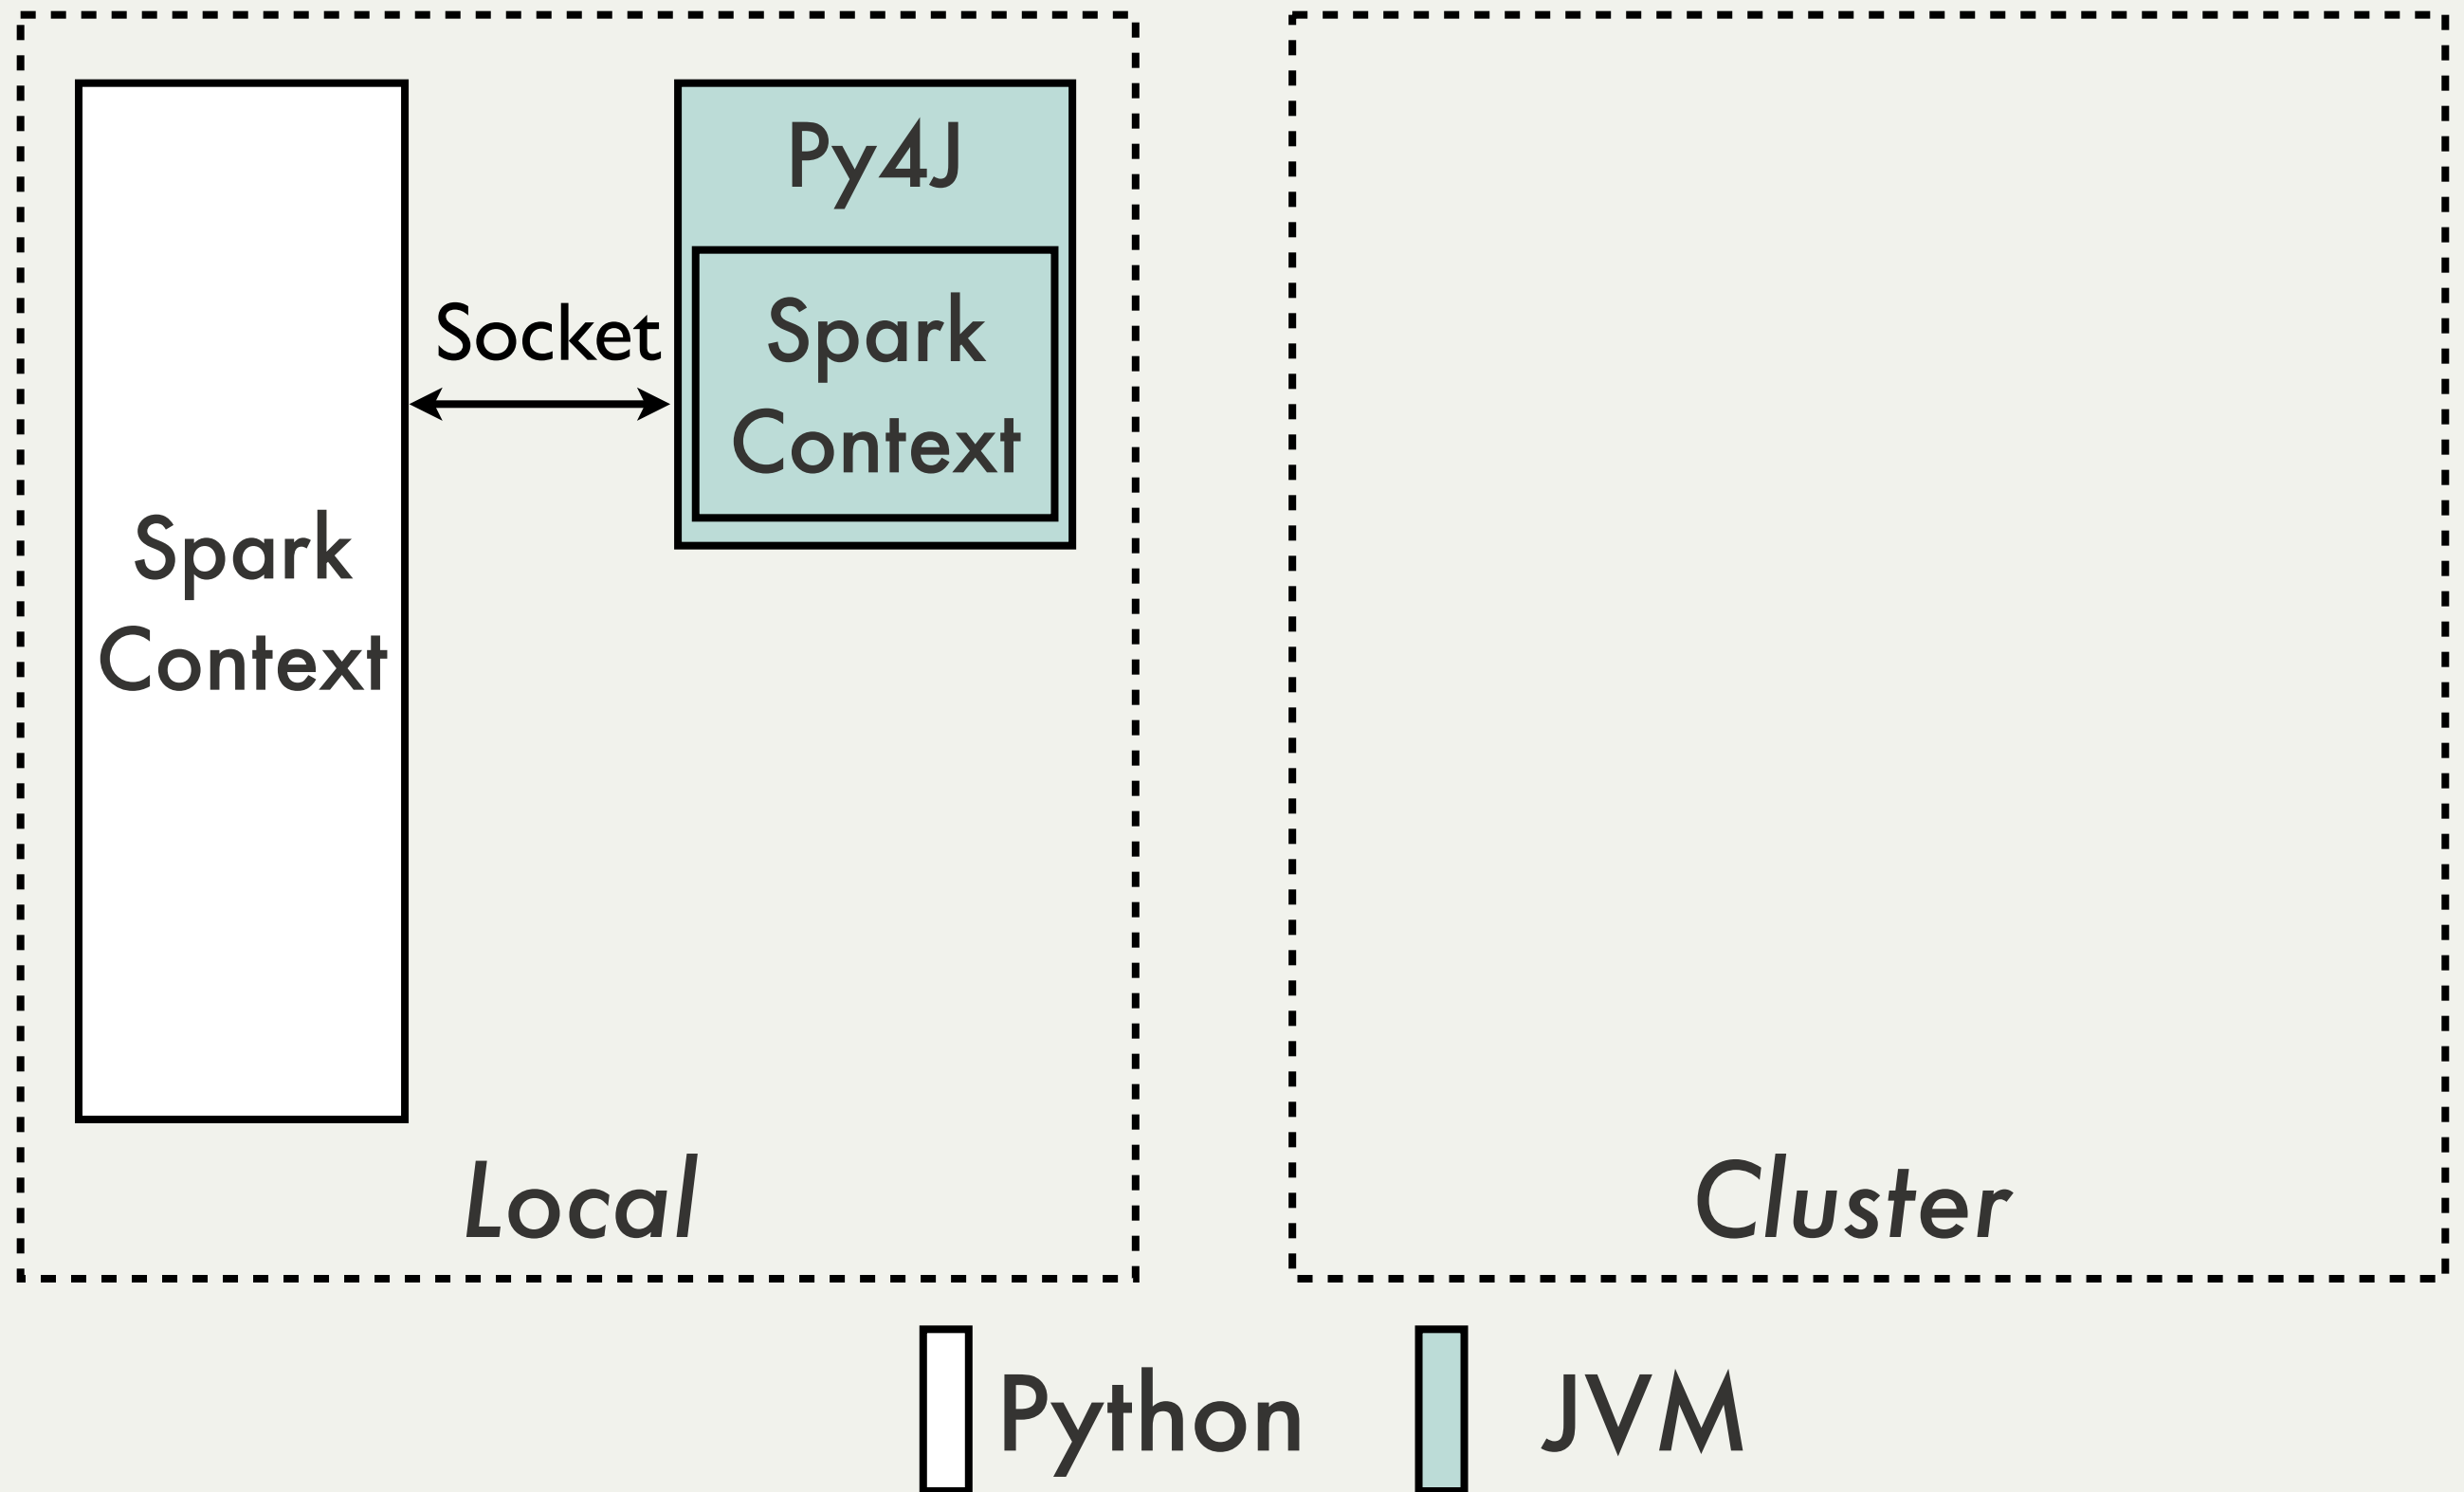
 Python

 JVM

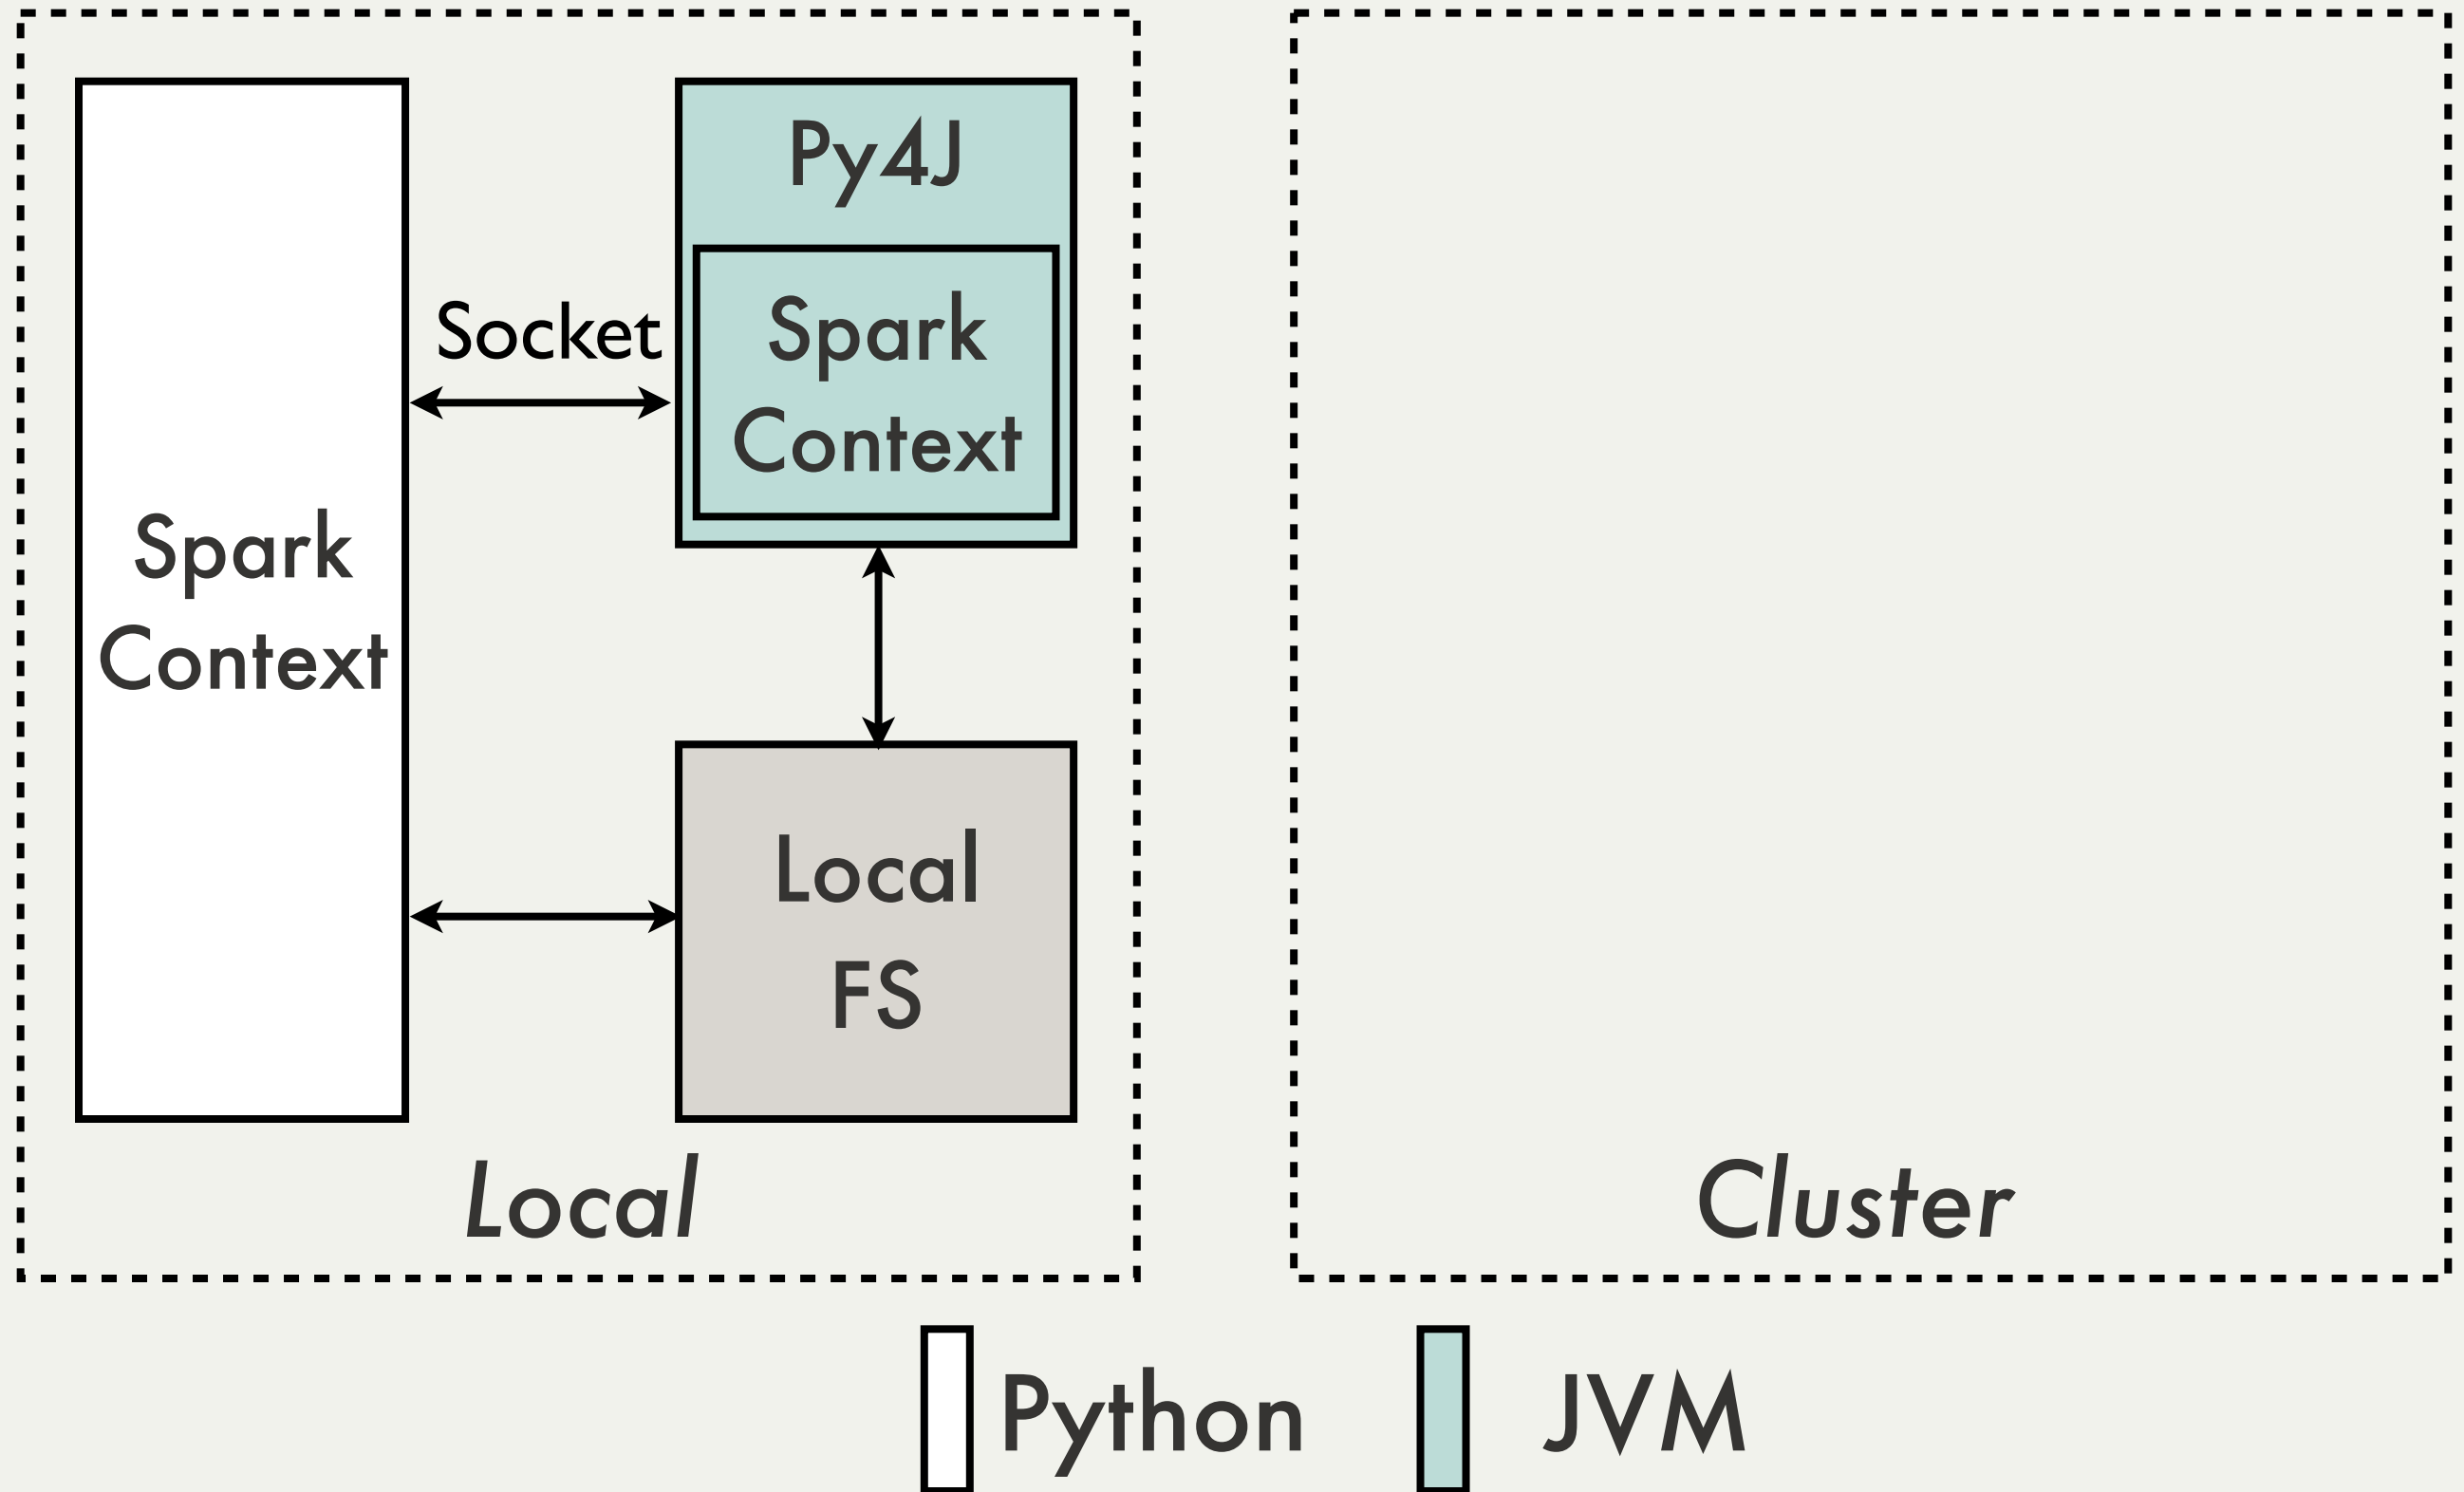
# Data Flow



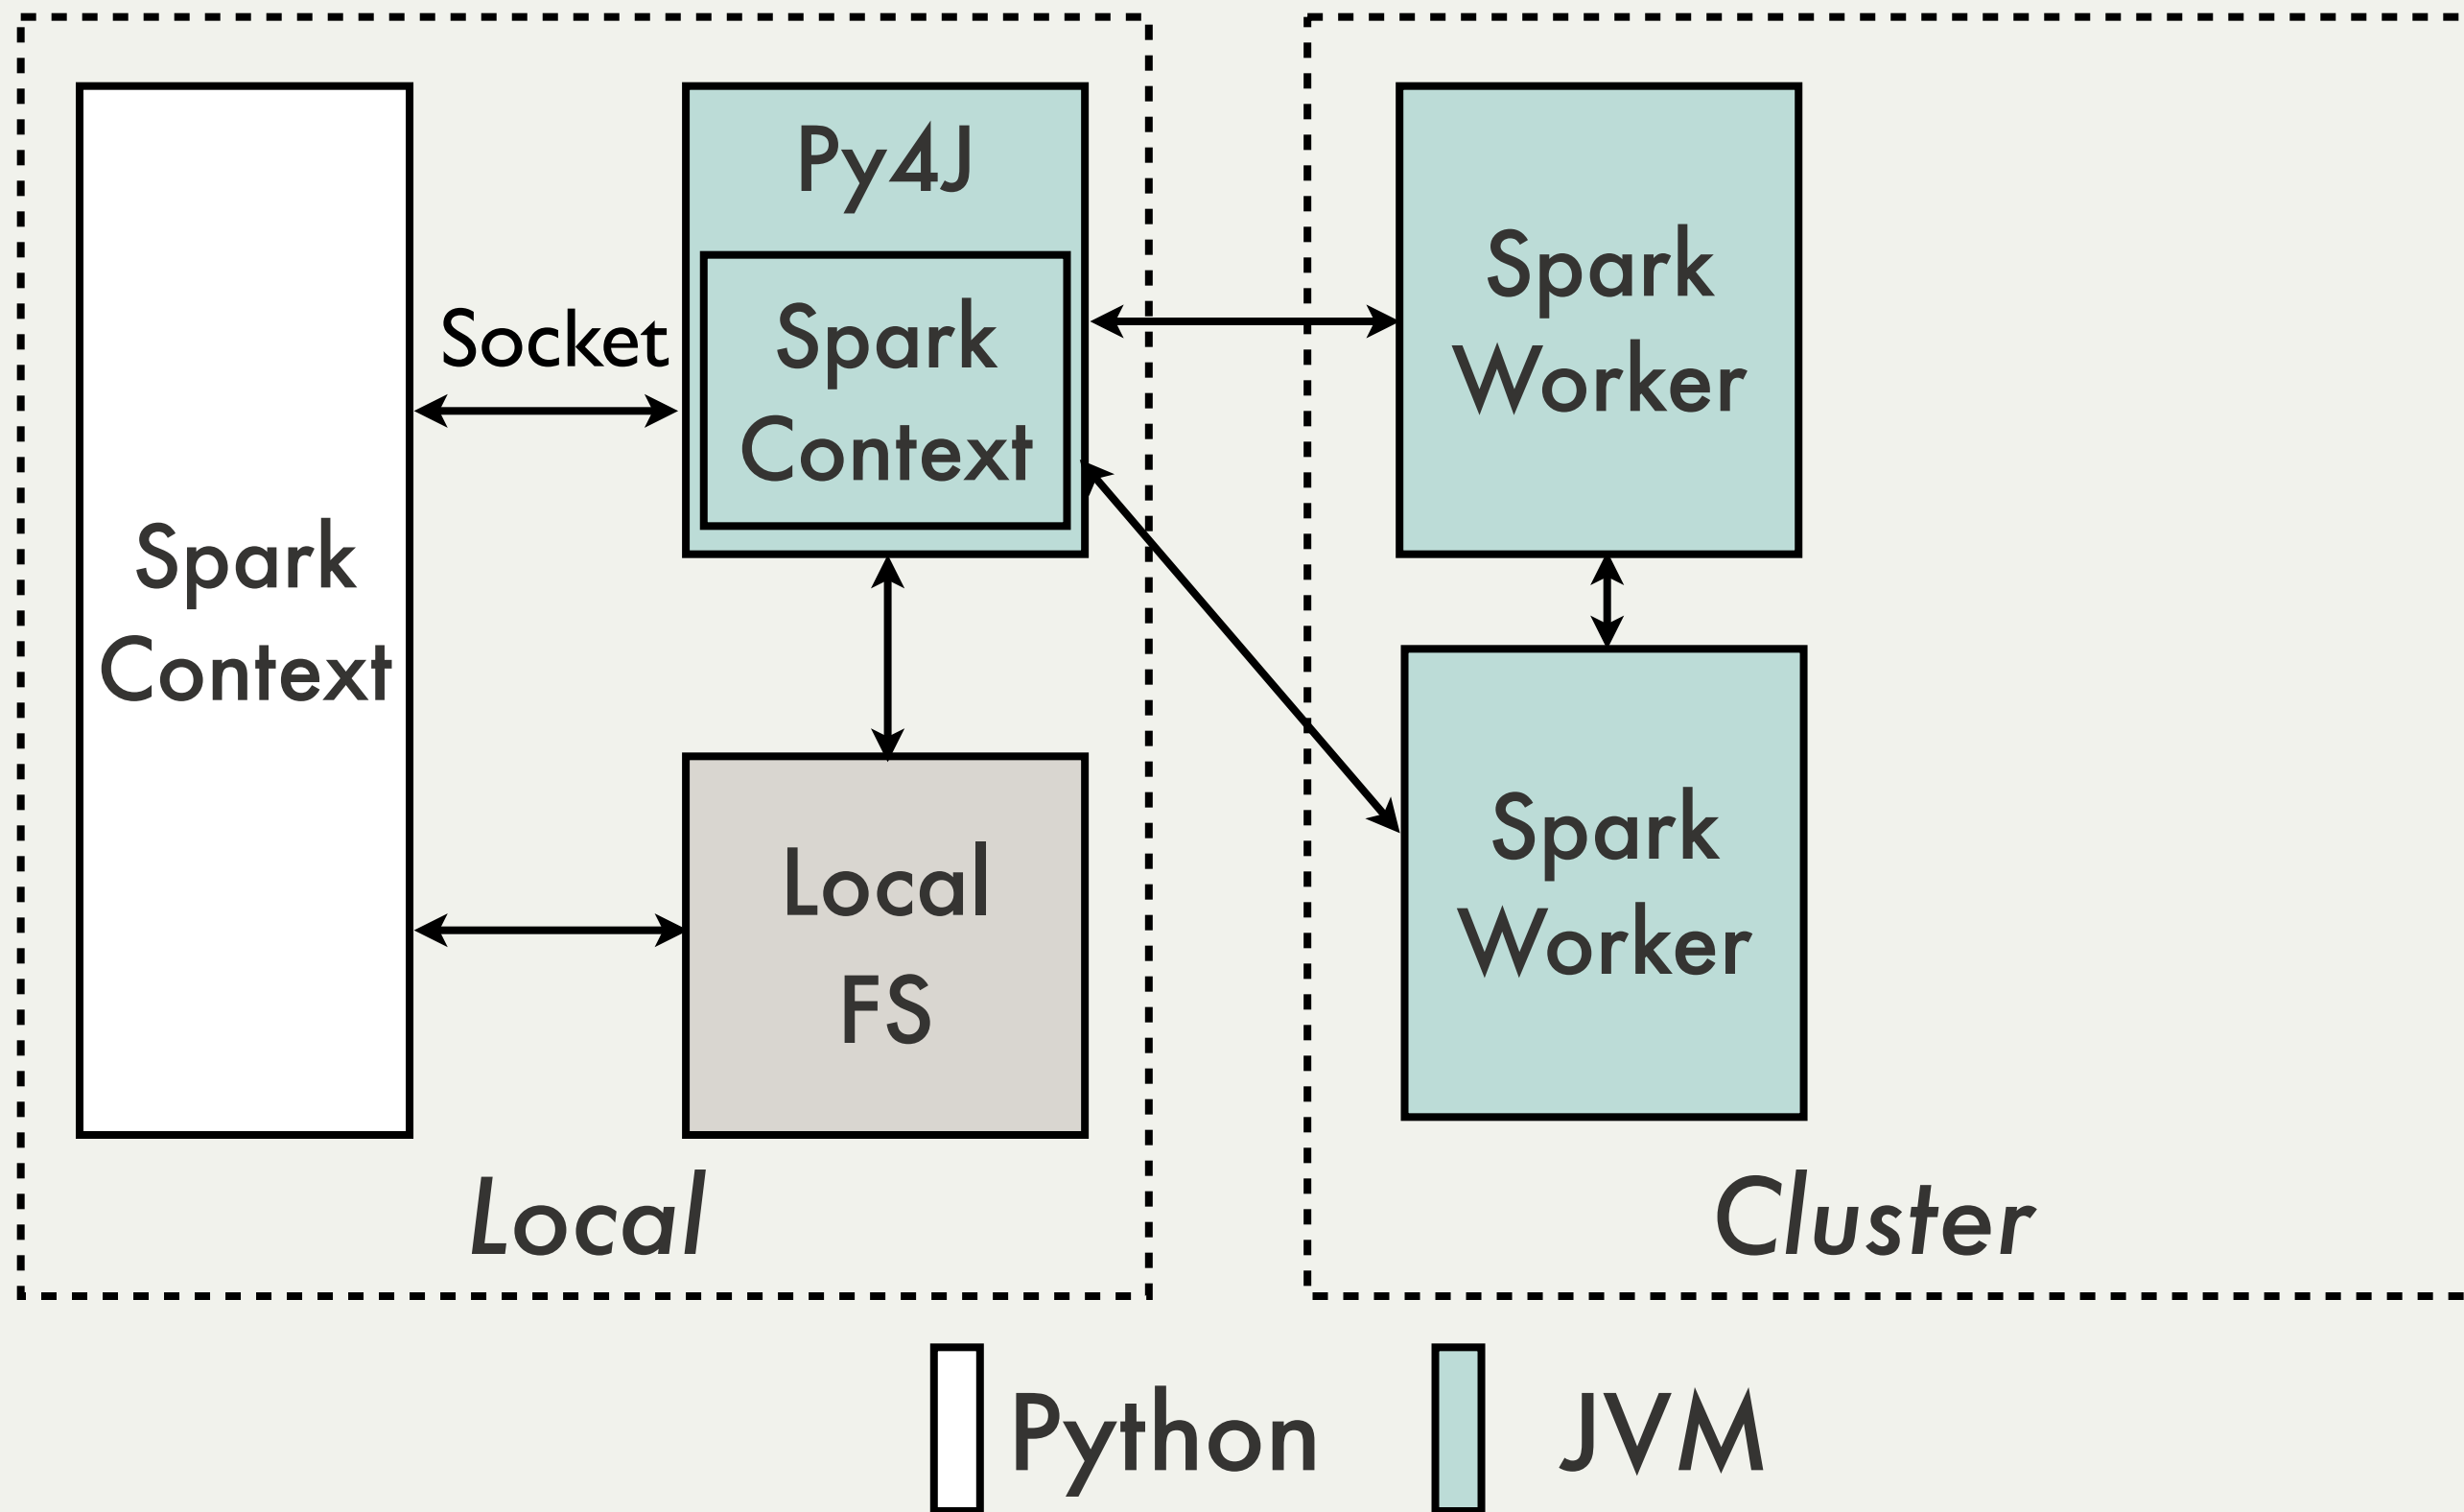
# Data Flow



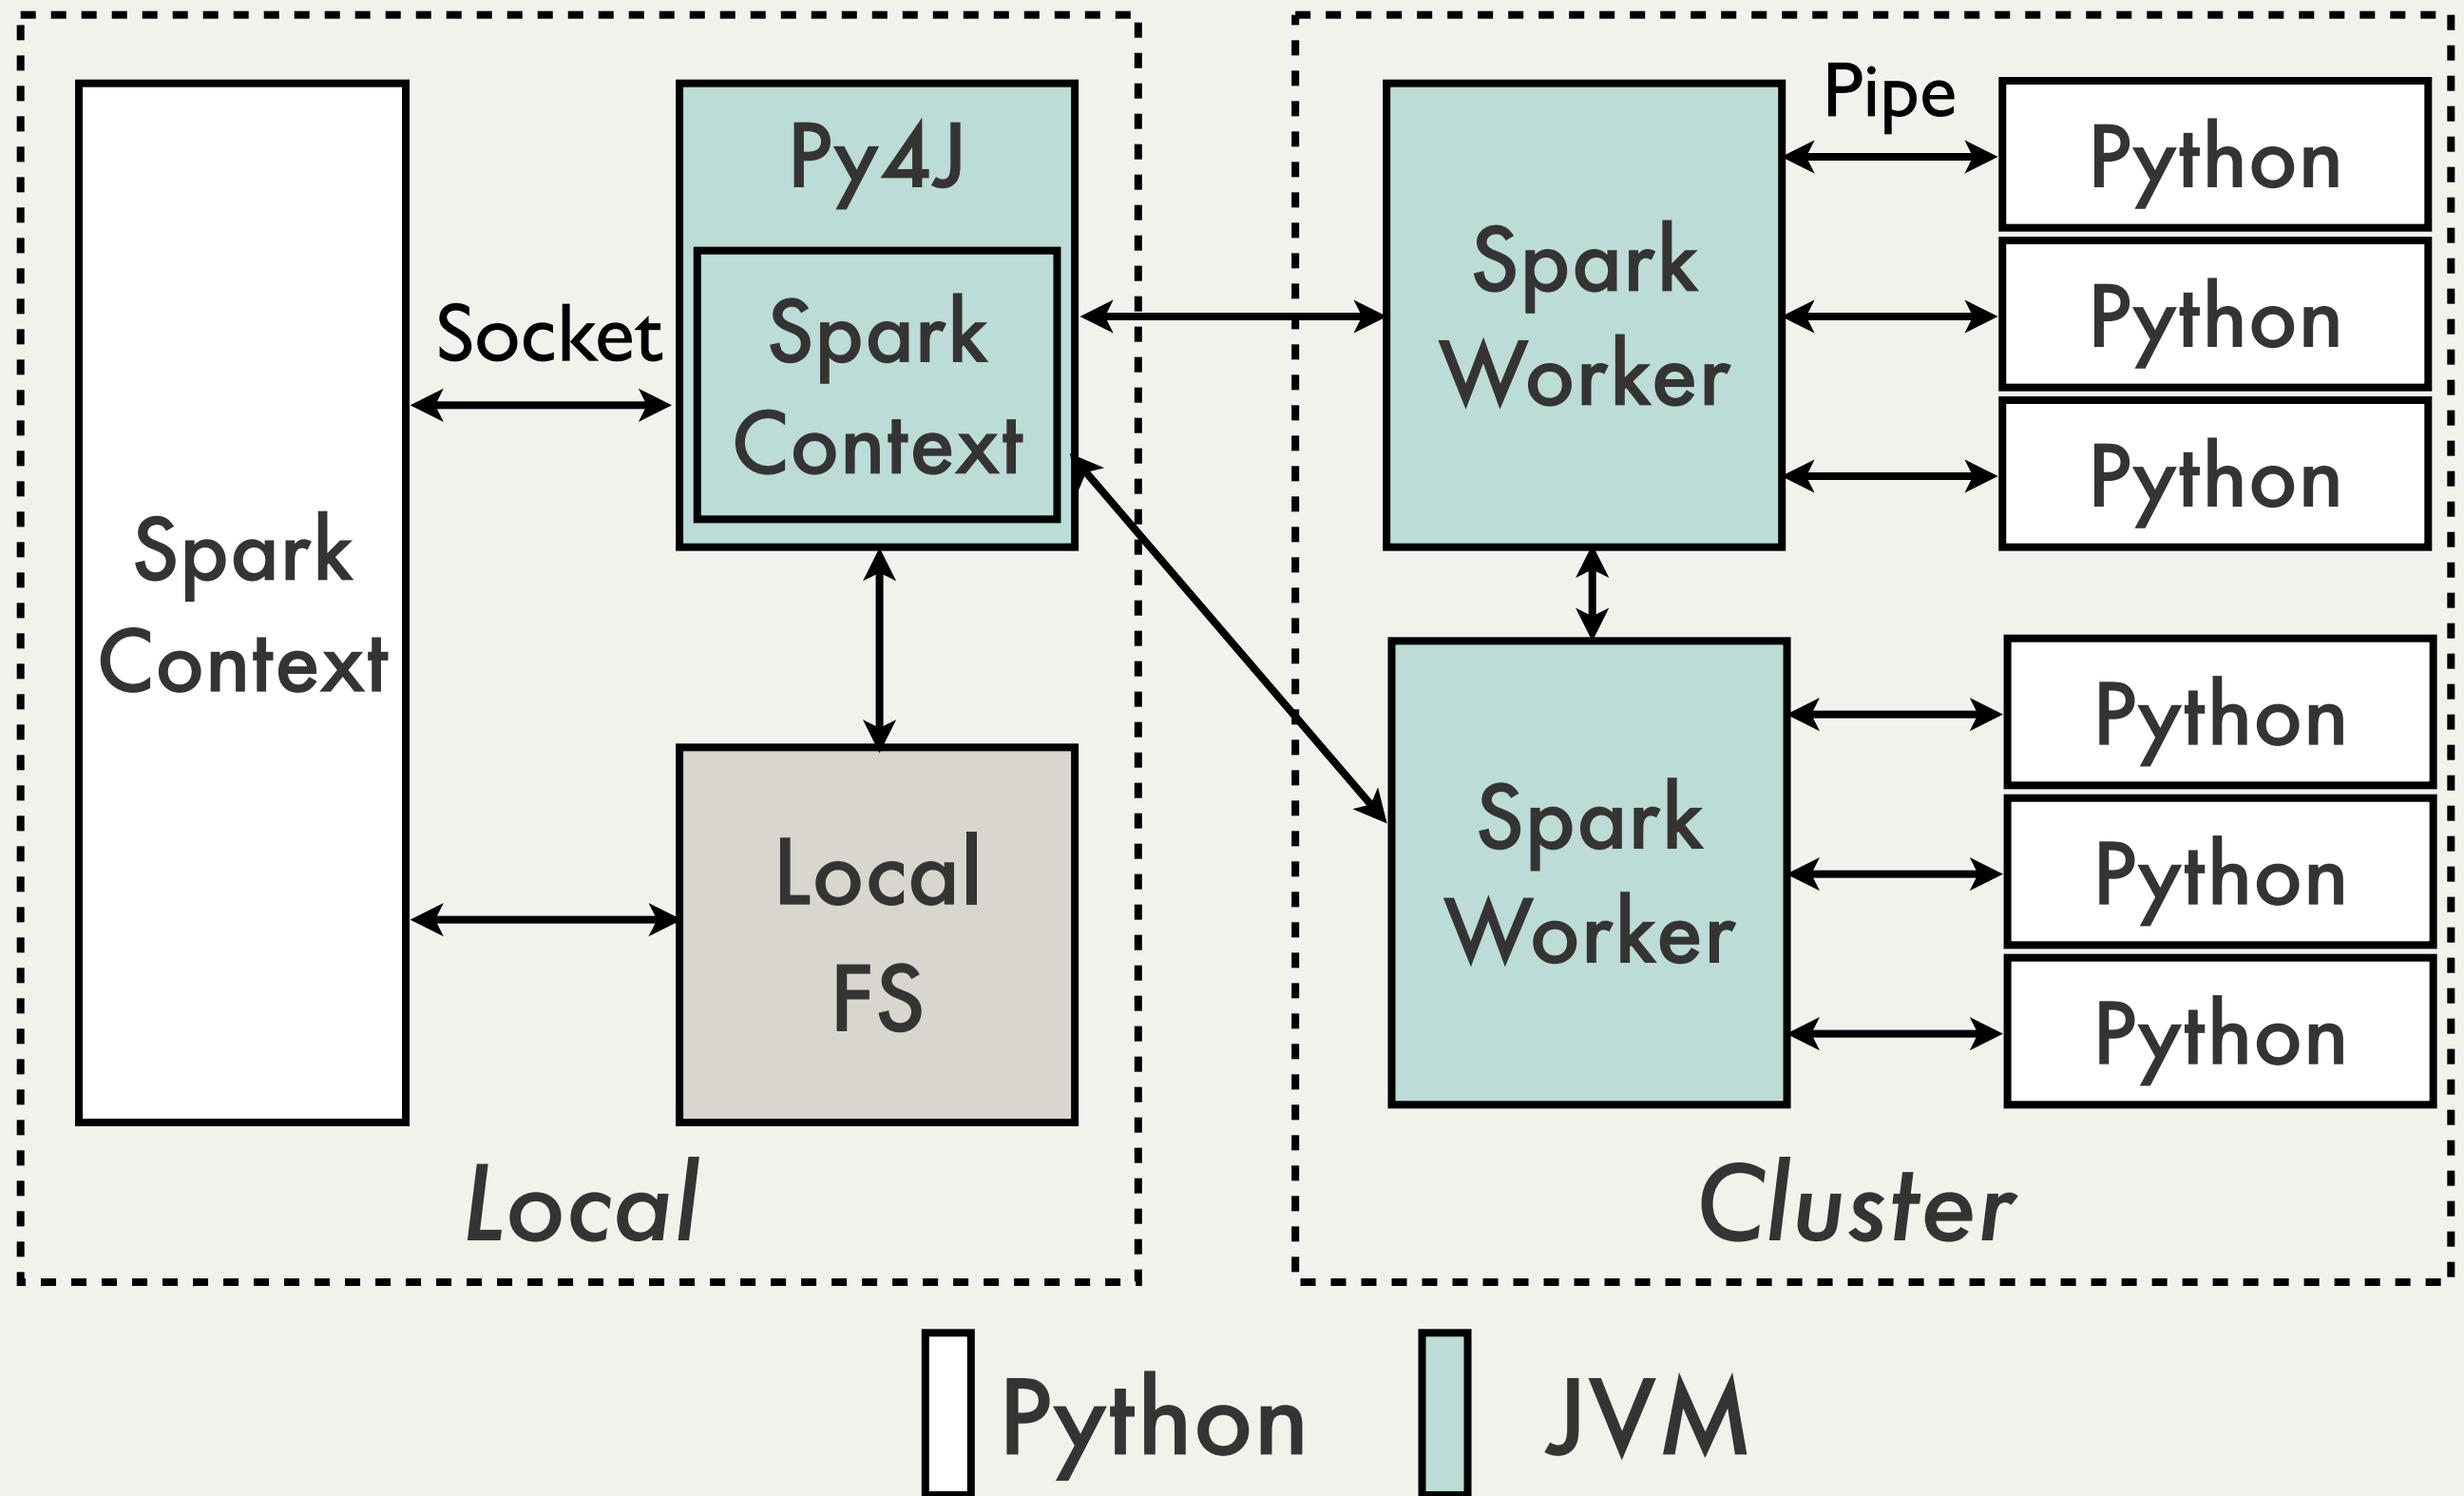
# Data Flow



# Data Flow



# Data Flow





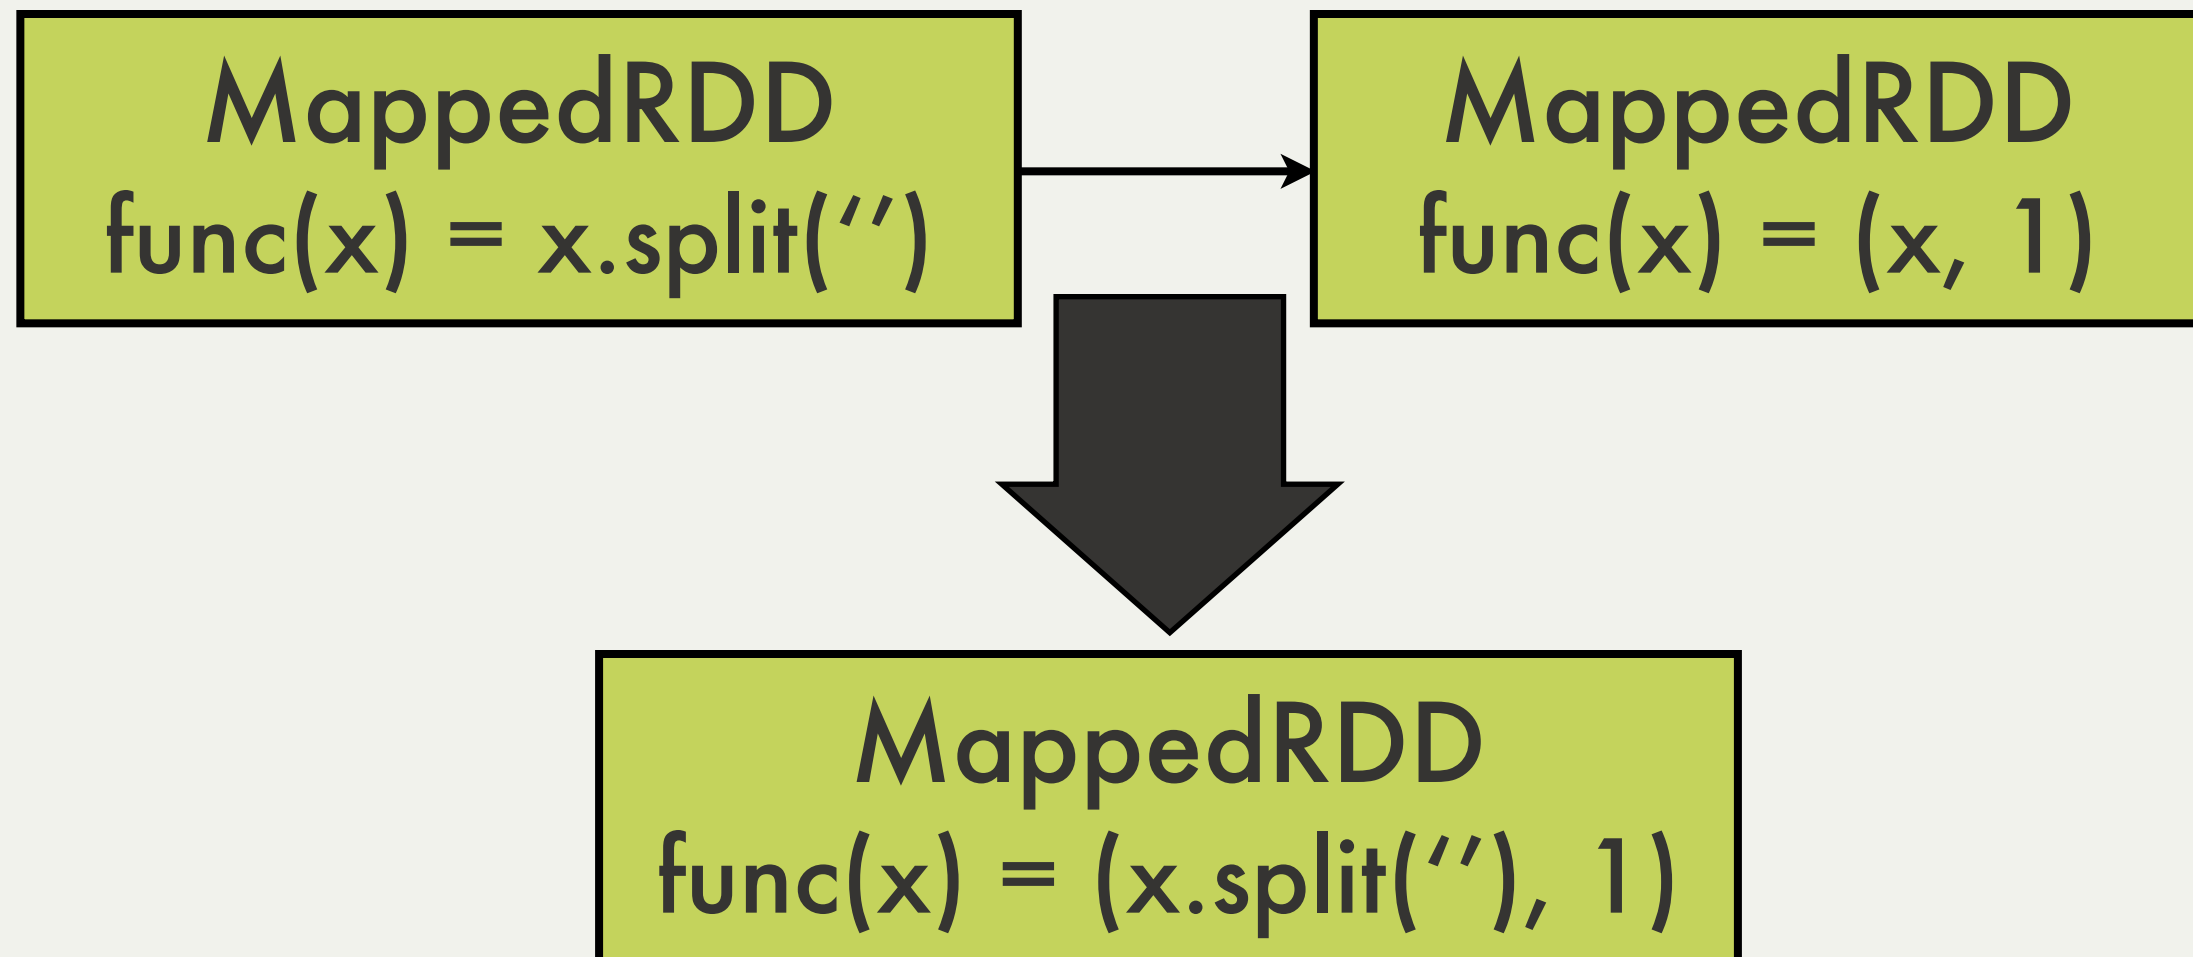


**Data is stored as Pickled  
objects in `RDD[Array[Byte]]`**

Storing *batches* of Python  
objects in one Scala object  
reduces overhead

# When possible, RDD transformations are pipelined:

```
lines.flatMap(lambda x: x.split(' ')) \
      .map(lambda x: (x, 1))
```



Python functions and closures  
are serialized using PiCloud's  
**CloudPickle** module

# Roadmap

Available in **Spark 0.7**



Thanks!

# Bonus Slides



# Pickle is a miniature **stack** **language**

```
>>> x = ["Hello", "World!"]
>>> pickletools.dis(cPickle.dumps(x, 2))
0: \x80  PROTO          2
2: ]      EMPTY_LIST
3: q      BININPUT      1
5: (      MARK
6: U      SHORT_BINSTRING 'Hello'
13: q     BININPUT      2
15: U     SHORT_BINSTRING 'World!'
23: q     BININPUT      3
25: e     APPENDS        (MARK at 5)
26: .     STOP
highest protocol among opcodes = 2
```

You can do crazy stuff, like  
converting a collection of  
pickled objects into a pickled  
collection.

<https://gist.github.com/JoshRosen/3384191>

# Bulk depickling can be faster *even if it involves Pickle opcode manipulation:*

10000 integers:

Bulk depickle (chunk size = 2): 0.266709804535

Bulk depickle (chunk size = 10): 0.0797798633575

Bulk depickle (chunk size = 100): 0.0388460159302

Bulk depickle (chunk size = 1000): 0.0333180427551

Individual depickle: 0.0540158748627

10000 dicts (dict([ (str(n), n) for n in range(100) ])):

Bulk depickle (chunk size = 2): 2.70617198944

Bulk depickle (chunk size = 10): 2.30310201645

Bulk depickle (chunk size = 100): 2.22087192535

Bulk depickle (chunk size = 1000): 2.22118020058

Individual depickle: 2.44124102592

<https://gist.github.com/JoshRosen/3401373>